

TechStack AI

7-MINUTE VIDEO PRESENTATION & LIVE DEMO SCRIPT

INSTRUCTIONS FOR RECORDER & SPEAKER: This script synchronizes on-screen visual capture actions with spoken narrative dialogue. Visual cues are boxed in light gray, and spoken dialogue blocks are highlighted in light blue.

Segment 1: Introduction & Platform Moat

0:00 - 1:00

ON-SCREEN ACTION (WHAT TO SHOW):

Show the main landing page (<http://localhost:5173>). Slowly scroll down to show the clean off-white background and branding highlight panels. Hover your cursor over the "Select Preset" buttons.

SPOKEN VOICEOVER DIALOGUE (WHAT TO SAY):

Welcome to TechStack AI—an Engineering Decision Platform designed to transform system design from a game of gut-feelings into a programmatically proven science. In modern software development, deciding on an architecture is often plagued by endless debates, outdated templates, and misaligned priorities. TechStack AI solves this. Rather than functioning as a generic chatbot, it executes a structured engineering decision workflow.

Segment 2: Live Operations & Stepper

1:00 - 2:30

ON-SCREEN ACTION (WHAT TO SHOW):

Click the "Uber Clone" preset button. Point to the Autofilled expected users and description fields, then click "Analyze System". Focus the camera on the live stepper activity progress spinners and streaming logs.

SPOKEN VOICEOVER DIALOGUE (WHAT TO SAY):

Let's see the platform in action. We'll load up a preset configuration—in this case, an Uber Clone—and trigger the analysis. Instantly, the platform launches a multi-stage background pipeline. Unlike traditional models, this is a sequence of dedicated AI agents working in harmony. First, the Requirements Analyst extracts constraints. Next, the Exa AI Research Agent performs a semantic search to retrieve engineering articles, followed by Systems, FinOps, and Risk agents. Each step is fully traced, streaming live logs directly to our screen.

Segment 3: Architecture Scorecard & Winner

2:30 - 3:30

ON-SCREEN ACTION (WHAT TO SHOW):

Scroll down to the "Candidate Architecture Leaderboard". Hover over cards for "Startup Optimized", "Balanced Scale", and "Enterprise Scale". Point to the final score and Winner label.

SPOKEN VOICEOVER DIALOGUE (WHAT TO SAY):

The pipeline is complete, and we are presented with our final scorecard. Instead of generating a single system recommendation, TechStack AI evaluates three distinct candidates: Startup Optimized, Balanced Scale, and Enterprise Scale. Every candidate is scored across four key engineering dimensions: Scalability, Cost, Reliability, and Development Speed. Crucially, the winning architecture is determined programmatically by our scoring engine in the codebase, ensuring complete transparency.

Segment 4: Benchmarks & Observability Traces

3:30 - 4:30

ON-SCREEN ACTION (WHAT TO SHOW):

Scroll down to show the "Benchmark Similarity" card showing progress meters, and point to the "Decision Trace" card showing stats (e.g. Sources sourced, papers analyzed).

SPOKEN VOICEOVER DIALOGUE (WHAT TO SAY):

To ground our recommendations, the engine compares the proposed designs against a local dataset of top engineering architectures, including Notion, Linear, Stripe, and Shopify. By running keyword-matching algorithms, the platform outputs a similarity percentage, showing you how closely your proposed stack aligns with established industry standards. Underneath, the Decision Trace panel displays raw observability stats—such as the exact number of requirements extracted, search citations processed, and benchmarks evaluated—proving the depth of the multi-agent workflow.

Segment 5: Interactive What-If Simulator

4:30 - 5:30

ON-SCREEN ACTION (WHAT TO SHOW):

Scroll to the "What-If Simulator". Drag the scale slider from 10k users, to 100k, and up to 1,000,000 users. Pause to let the leaderboard scores and cost growth curves update on the client in real-time.

SPOKEN VOICEOVER DIALOGUE (WHAT TO SAY):

One of the most powerful features of the platform is the interactive What-If Simulator. Using this slider, we can simulate traffic growth in real-time. Let's drag the scale up to one million users. Watch how the dashboard reacts. The scoring engine dynamically re-adjusts its priorities. At low scale, development speed and cost are prioritized. At high scale, scalability and reliability take over. As a result, the scores recalculate instantly, selecting a new winner and projecting new growth cost curves—all locally on the client.

Segment 6: Backend Architecture Tour

5:30 - 6:30

ON-SCREEN ACTION (WHAT TO SHOW):

Switch to IDE. Show routes (`analyzeRoutes.ts`), controller (`analyzeController.ts`), paced scheduler (`orchestratorService.ts`), and scoring engine (`scoring.ts`).

SPOKEN VOICEOVER DIALOGUE (WHAT TO SAY):

Let's take a quick look behind the scenes at the backend architecture. Our router defines clean REST endpoints for analyze, reports, and health checking. When an analysis request is submitted, the controller validates inputs and hands execution off to the Orchestrator. The Orchestrator manages the sequential activation of our agents, introducing deliberate pacing delays to respect API rate limits. Finally, the scoring engine utility calculates the scores using dynamic weighting matrices based on expected user traffic tiers, ensuring a deterministic decision-making process.

Segment 7: Outro & Summary

6:30 - 7:00

ON-SCREEN ACTION (WHAT TO SHOW):

Switch back to the browser showing the database schema tables at the bottom of the page, then scroll back to the top header.

SPOKEN VOICEOVER DIALOGUE (WHAT TO SAY):

By combining multi-agent research pipelines, programmatic scoring models, and interactive simulation calculators, TechStack AI delivers a truly impressive Engineering Decision Platform. It bridges the gap between high-level ideas and production-grade architectures. Thank you for watching this demonstration.

