

TechStack AI

7-MINUTE DEMO SCRIPT - PLAIN ENGLISH VERSION

INSTRUCTIONS FOR THE VOICE ACTOR: Read only the blue-highlighted boxes. All coding terms, file extensions (like .ts), and technical jargon have been replaced with clear, natural spoken English.

Segment 1: Introduction & Platform Moat

0:00 - 1:00

ON-SCREEN ACTION (WHAT TO SHOW):

Show the main landing page (<http://localhost:5173>). Slowly scroll down to show the clean off-white background and branding highlight panels. Hover your cursor over the "Select Preset" buttons.

SPOKEN VOICE LINES (WHAT TO SAY - PLAIN ENGLISH):

Welcome to TechStack AI—an Engineering Decision Platform designed to transform system design from a game of gut-feelings into a programmatically proven science. In modern software development, deciding on an architecture is often plagued by endless debates, outdated templates, and misaligned priorities. TechStack AI solves this. Rather than functioning as a generic messaging window, it executes a structured engineering decision workflow.

Segment 2: Live Operations & Stepper

1:00 - 2:30

ON-SCREEN ACTION (WHAT TO SHOW):

Click the "Uber Clone" preset button. Point to the Autofilled expected users and description fields, then click "Analyze System". Focus the camera on the live stepper activity progress spinners and streaming logs.

SPOKEN VOICE LINES (WHAT TO SAY - PLAIN ENGLISH):

Let's see the platform in action. We'll load up a preset configuration—in this case, an Uber Clone—and trigger the analysis. Instantly, the platform launches a multi-stage background pipeline. Unlike traditional tools, this uses a team of dedicated digital assistants working in harmony. First, the requirements assistant extracts constraints. Next, the web search assistant performs a search to retrieve real engineering articles, followed by systems, costing, and safety auditors. Each step is fully tracked, streaming live logs directly to our screen.

Segment 3: Architecture Scorecard & Winner

2:30 - 3:30

ON-SCREEN ACTION (WHAT TO SHOW):

Scroll down to the "Candidate Architecture Leaderboard". Hover over cards for "Startup Optimized", "Balanced Scale", and "Enterprise Scale". Point to the final score and Winner label.

SPOKEN VOICE LINES (WHAT TO SAY - PLAIN ENGLISH):

The analysis is complete, and we are presented with our final scorecard. Instead of generating a single system recommendation, TechStack AI evaluates three distinct candidates: Startup Optimized, Balanced Scale, and Enterprise Scale. Every candidate is scored across four key dimensions: Scalability, Cost, Reliability, and Development Speed. Crucially, the winning architecture is determined programmatically by our scoring engine in the codebase, ensuring complete transparency.

Segment 4: Benchmarks & Observability Traces

3:30 - 4:30

ON-SCREEN ACTION (WHAT TO SHOW):

Scroll down to show the "Benchmark Similarity" card showing progress meters, and point to the "Decision Trace" card showing stats (e.g. Sources sourced, papers analyzed).

SPOKEN VOICE LINES (WHAT TO SAY - PLAIN ENGLISH):

To ground our recommendations, the engine compares the proposed designs against a local dataset of top engineering architectures, including Notion, Linear, Stripe, and Shopify. By running matching rules, the platform outputs a similarity percentage, showing you how closely your proposed stack aligns with established industry standards. Underneath, the observability panel displays raw statistics—such as the exact number of requirements extracted, search citations processed, and benchmarks evaluated—proving the depth of the automated workflow.

Segment 5: Interactive What-If Simulator

4:30 - 5:30

ON-SCREEN ACTION (WHAT TO SHOW):

Scroll to the "What-If Simulator". Drag the scale slider from 10k users, to 100k, and up to 1,000,000 users. Pause to let the leaderboard scores and cost growth curves update on the client in real-time.

SPOKEN VOICE LINES (WHAT TO SAY - PLAIN ENGLISH):

One of the most powerful features of the platform is the interactive traffic simulator. Using this slider, we can simulate traffic growth in real-time. Let's drag the scale up to one million users. Watch how the dashboard reacts. The scoring calculator dynamically adjusts its priorities. At low scale, development speed and cost are prioritized. At high scale, scalability and reliability take over. As a result, the scores recalculate instantly, selecting a new winner and projecting new growth cost curves—all locally on the browser.

Segment 6: Backend Architecture Tour

5:30 - 6:30

ON-SCREEN ACTION (WHAT TO SHOW):

Switch to IDE. Show routes (`analyzeRoutes.ts`), controller (`analyzeController.ts`), paced scheduler (`orchestratorService.ts`), and scoring engine (`scoring.ts`).

SPOKEN VOICE LINES (WHAT TO SAY - PLAIN ENGLISH):

Let's take a quick look behind the scenes at the server. Here, the server router file directs requests for analyzing, retrieving reports, and checking systems health. When an analysis request is submitted, the server controller file validates the inputs and hands them over to the agent coordinator file. The coordinator manages the team of digital assistants one by one, adding minor pauses to avoid overloading the systems. Finally, the scoring calculator file determines the final points using mathematical rules based on expected traffic.

Segment 7: Outro & Summary

6:30 - 7:00

ON-SCREEN ACTION (WHAT TO SHOW):

Switch back to the browser showing the database schema tables at the bottom of the page, then scroll back to the top header.

SPOKEN VOICE LINES (WHAT TO SAY - PLAIN ENGLISH):

By combining automated research, math scoring, and interactive traffic simulations, TechStack AI delivers a truly impressive Engineering Decision Platform. It bridges the gap between high-level ideas and production-grade architectures. Thank you for watching this demonstration.

